

Self-driving car engineer nano-degree

Project 5: Vehicle detection

The goals/steps of this project are the following:

- Perform a Histogram of Oriented Gradients (HOG) feature extraction on a labeled training set of images and train a classifier Linear SVM classifier.
- Optionally, apply a color transform and append binned color features, as well as histograms of color, to HOG feature vector.
- Implement a sliding-window technique and use the trained classifier to search for vehicles in images.
- Run the pipeline on a video stream and create a heat map of recurring detections frame by frame to reject outliers and follow detected vehicles.
- Estimate a bounding box for vehicles detected.

Here I will consider the rubric points individually and describe how I addressed each point in my implementation.

In order for the steps to be successful I first import all the required libraries that I will make use of. Then I define a function *get_hog_features*, which takes in:

- An image, on which HOG features will be extracted .
- Orient, the number of orientation bins that the gradient information will be split up into in the histogram.

- `pix_per_cell`, which specifies the cell size over which each gradient histogram is computed.
- `cell_per_block`, the local area over which the histogram counts in a given cell will be normalized.
- `vis` which is a Boolean in order to determine whether or not the function should return the hog image along with the features
- `Feature_vec`, a Boolean to automatically unroll the features

Next I define a function called *bin_spatial* which computes the binned color features given an image and a new image size.

I then define a function called *color_hist* which computes the color histogram features on each color channel separately and combines the histograms into a single feature vector.


The three previously mentioned functions is called and utilized through the *extract_features* function which is used to extract features from a list of images. This function takes in:

- A list of images
- The desired color space
- The spatial size required for the *bin_spatial* function
- The number of histogram bins
- Orient
- `pix_per_cell`
- `cell_per_block`
- hog channel
- `spatial_feat`, Used to determine whether or not to use the *bin_spatial* function.

- hist_feat, used to determine whether or not to use the color_hist function.
- hog_feat, used to determine whether or not to use the get_hog_features function.

This function then returns a list of feature vectors.

Next I load all the images required for training the classifier into the cars and notcars arrays. And then display a randomly selected image from each array:




After displaying an image from each array, I then define the parameters that I have chosen for the *extract_features* function:

- ❖ color_space: YCrCb
- ❖ orient: 11
- ❖ pix_per_cell: 16
- ❖ cell_per_block: 2
- ❖ hog_channel: ALL
- ❖ special_size: (32, 32)
- ❖ hist_bins: 32

- ❖ spatial_feat: True
- ❖ hist_feat: True
- ❖ hog_feat: True
- ❖ y_start_stop: [380, 656], The minimum and maximum in y to search in slide window.

After defining the parameters I use them to extract features from the cars and notcars arrays into car_features and notcar_features. I then convert a randomly selected car image to grayscale and extract hog features from the image and then display the original along with the HOG features:



After the feature extraction I train an SVC classifier. In order to achieve this I first create an array stack of feature vectors and then fit them to the **StandardScaler** function.

I then apply the scaler to the array stack feature vectors and define the labels for the vectors.

I then split the data into randomized training and test sets with a 80-20% training-validation split. I then fit the newly split and

randomized data on to the svc classifier and print/display the test accuracy obtained from the trained classifier:

```
Test Accuracy of SVC = 0.999
```

Next I define the *single_img_features* function which is similar to the *extreact_features* function except that the new function only takes in a single image.

Then I defined a *slide_window* function which returns a list of windows.

Next the *draw_boxes* function iterates through bounding boxes and draws a rectangle to on the given coordinates, and then returns the image containing the newly drawn rectangles.

The *search_windows* function searches through the windows provided by the *slide_window* function utilizing the *single_img_features* function. The extracted features is then fed into the classifier and performs a prediction on the given features within the windows. Then if the prediction yield a result of “1” the specified window is then added to the *on_windows* array which then contains all the windows predicted to be cars.

Next I loop though all the images contained within the *test_images* folder to determine whether the pipeline up to this point is capable of detecting cars within each image that a car is present:



After testing the pipeline on the test images I found that it detects the cars relatively well except that in a few places it detects cars where there are none. In order to address this problem I created three new functions:

- `add_heat`:
 - This function iterates through a list of bounding boxes, and adds “1” for all pixels inside the bounding box. Then returns the updated heat map.
- `apply_threshold`:
 - This function zeros out all pixels below the selected threshold and returns a thresholded map.
- `draw_labeled_bboxes`:
 - This function iterates through all the detected cars from the previous functions and draws a rectangle determined through the labels value.

Next I define the `vid_to_img` function which utilizes the previously created pipeline. Here once again I use all the images stored in the `test_images` folder to determine whether the newly created functions improve the vehicle detection pipeline:



The resulting images shows that the new functions did indeed remove the previously false detected windows.

I then proceeded to call the *vid_to_img* function on the `project_video.mp4` which is saved in the output_video folder.

Conclusion

As seen by viewing the `project_video.mp4` in the `output_video` folder the pipeline is completely capable of detecting cars within a video stream. However the pipeline does seem to struggle detecting the white cars for a couple of seconds possibly due to the car being out of range due to the `y_start_stop` value. Other possible improvements to the pipeline could be to increase the speed at which the pipeline performs on the video stream as it is currently taking very long to process.